

Caldo de Kafka

(para Node)

Alberto Avila
albertoavila@gmail.com
@alberteiu



Disclaimer

Not an actual javascript talk (sorry!)



What is Kafka?

- Pub/Sub messaging rethought as a distributed commit log.

The log¹ 🌲

¹ The Log: What every software engineer should know about real-time data's unifying abstraction <http://bit.do/thelog>

Append only ordered sequence of records

- Not unstructured error messages (i.e dmesg)
- Sequence of inputs (events) on a system
- Useful for distributed systems



How to replicate my application state

1. On every request you write a log entry with the request itself and update your state
2. Replicas read from the log and applies the logged request to its current state

💣 **Hey, you just (re)invented DB replication!** 💣



Let me sidetrack you for a moment

Real world problems

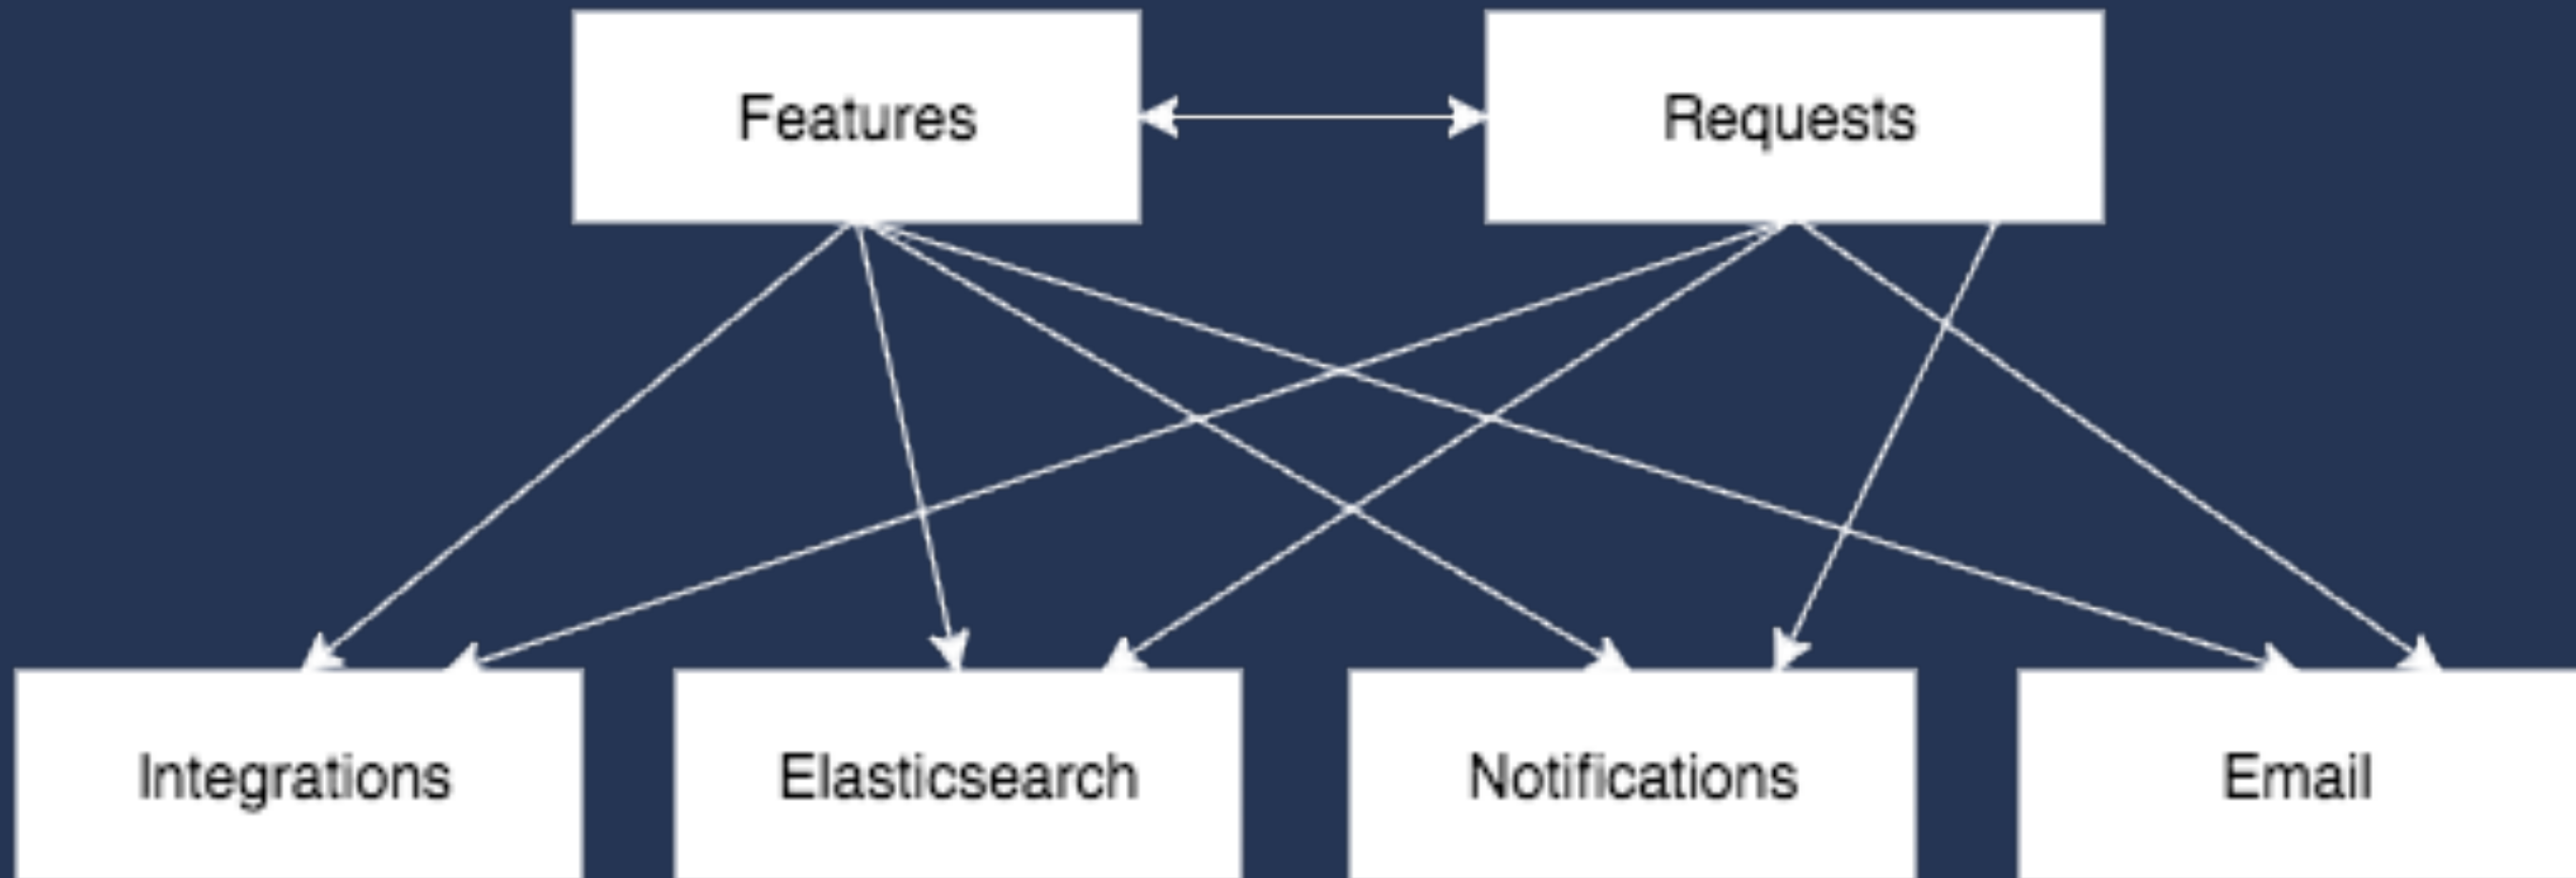
Meet Wizeline

A user creates a Feature

- Now you need to send an email
- And a push notification
- And index it on Elasticsearch
- And also need to update the information on Jira
- And salesforce too!

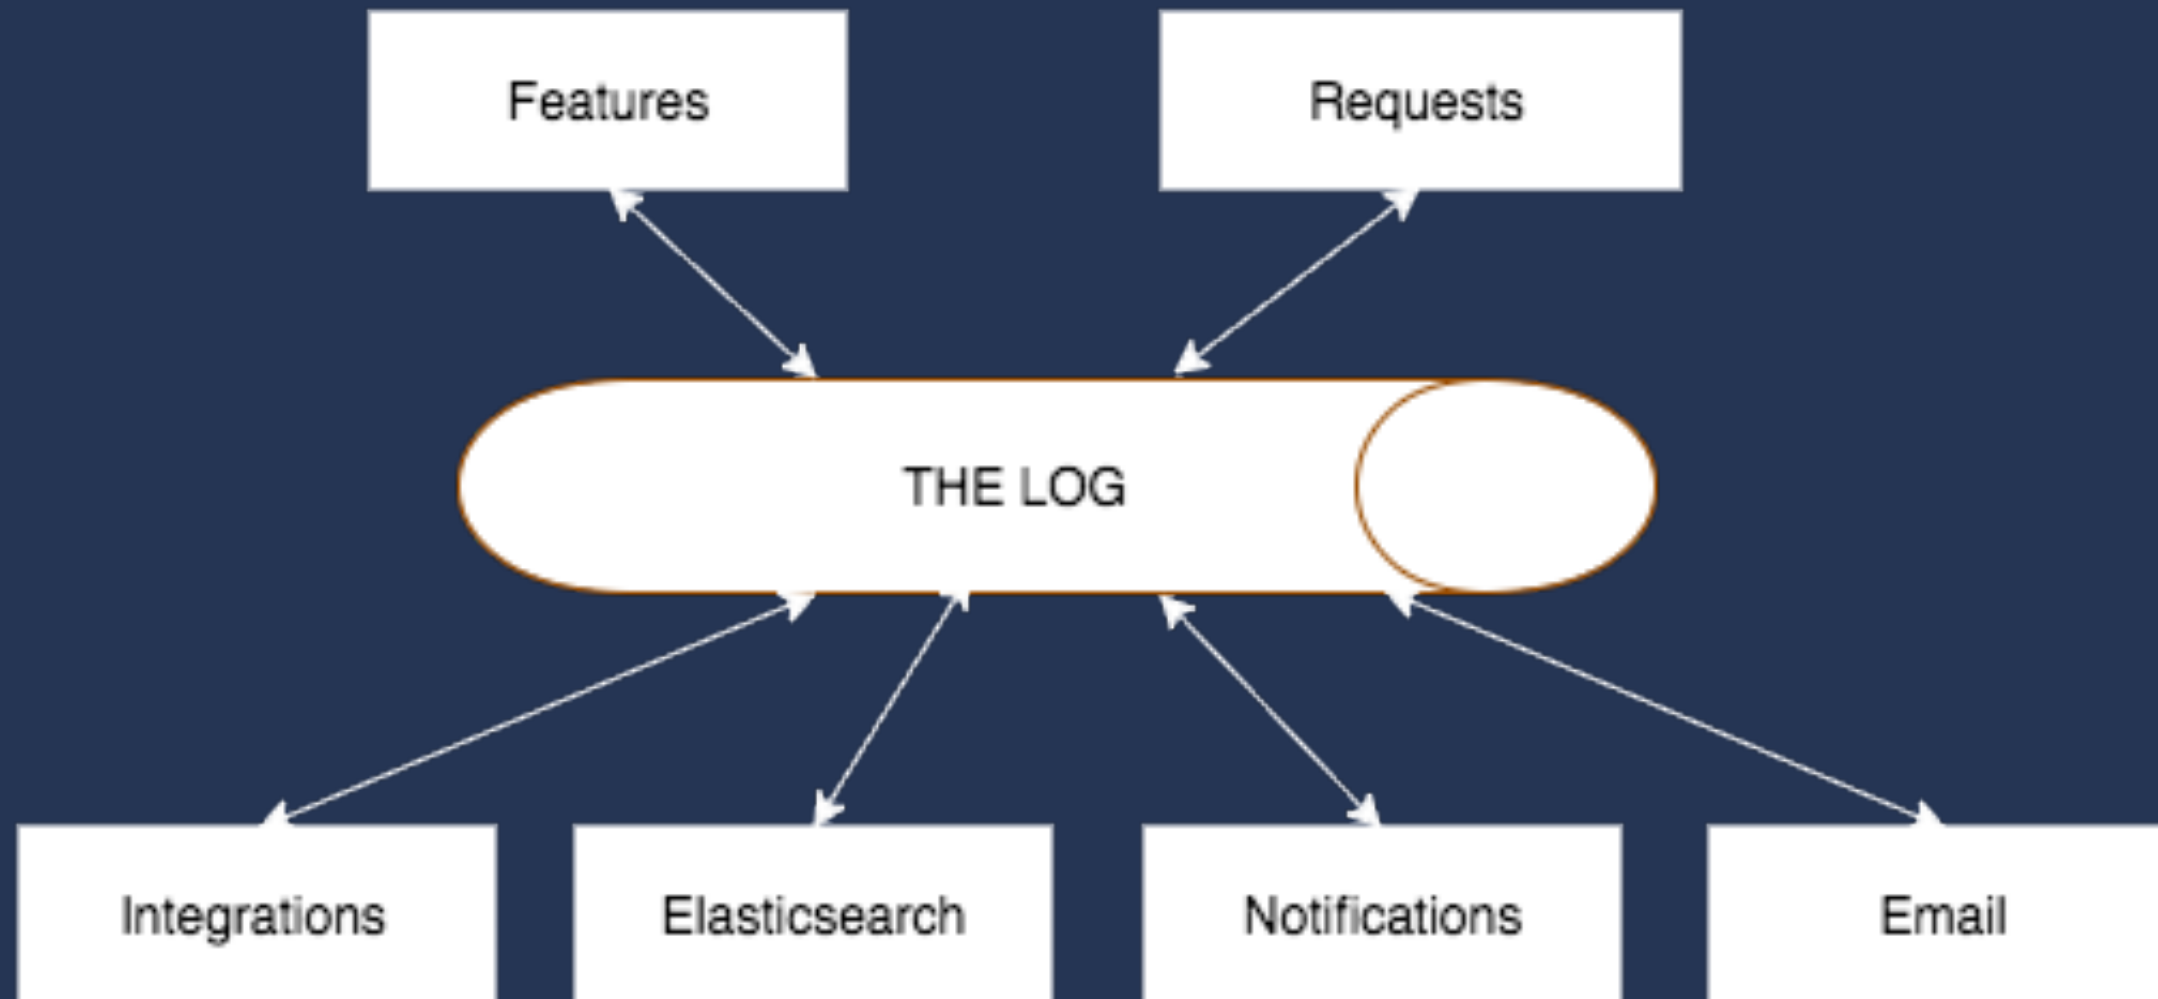
Notice a trend?

n **systems** $\rightarrow n^2$ **connections**



Yup, that was beyond my mouse-drawing skills.

Decoupling services through a centralized log

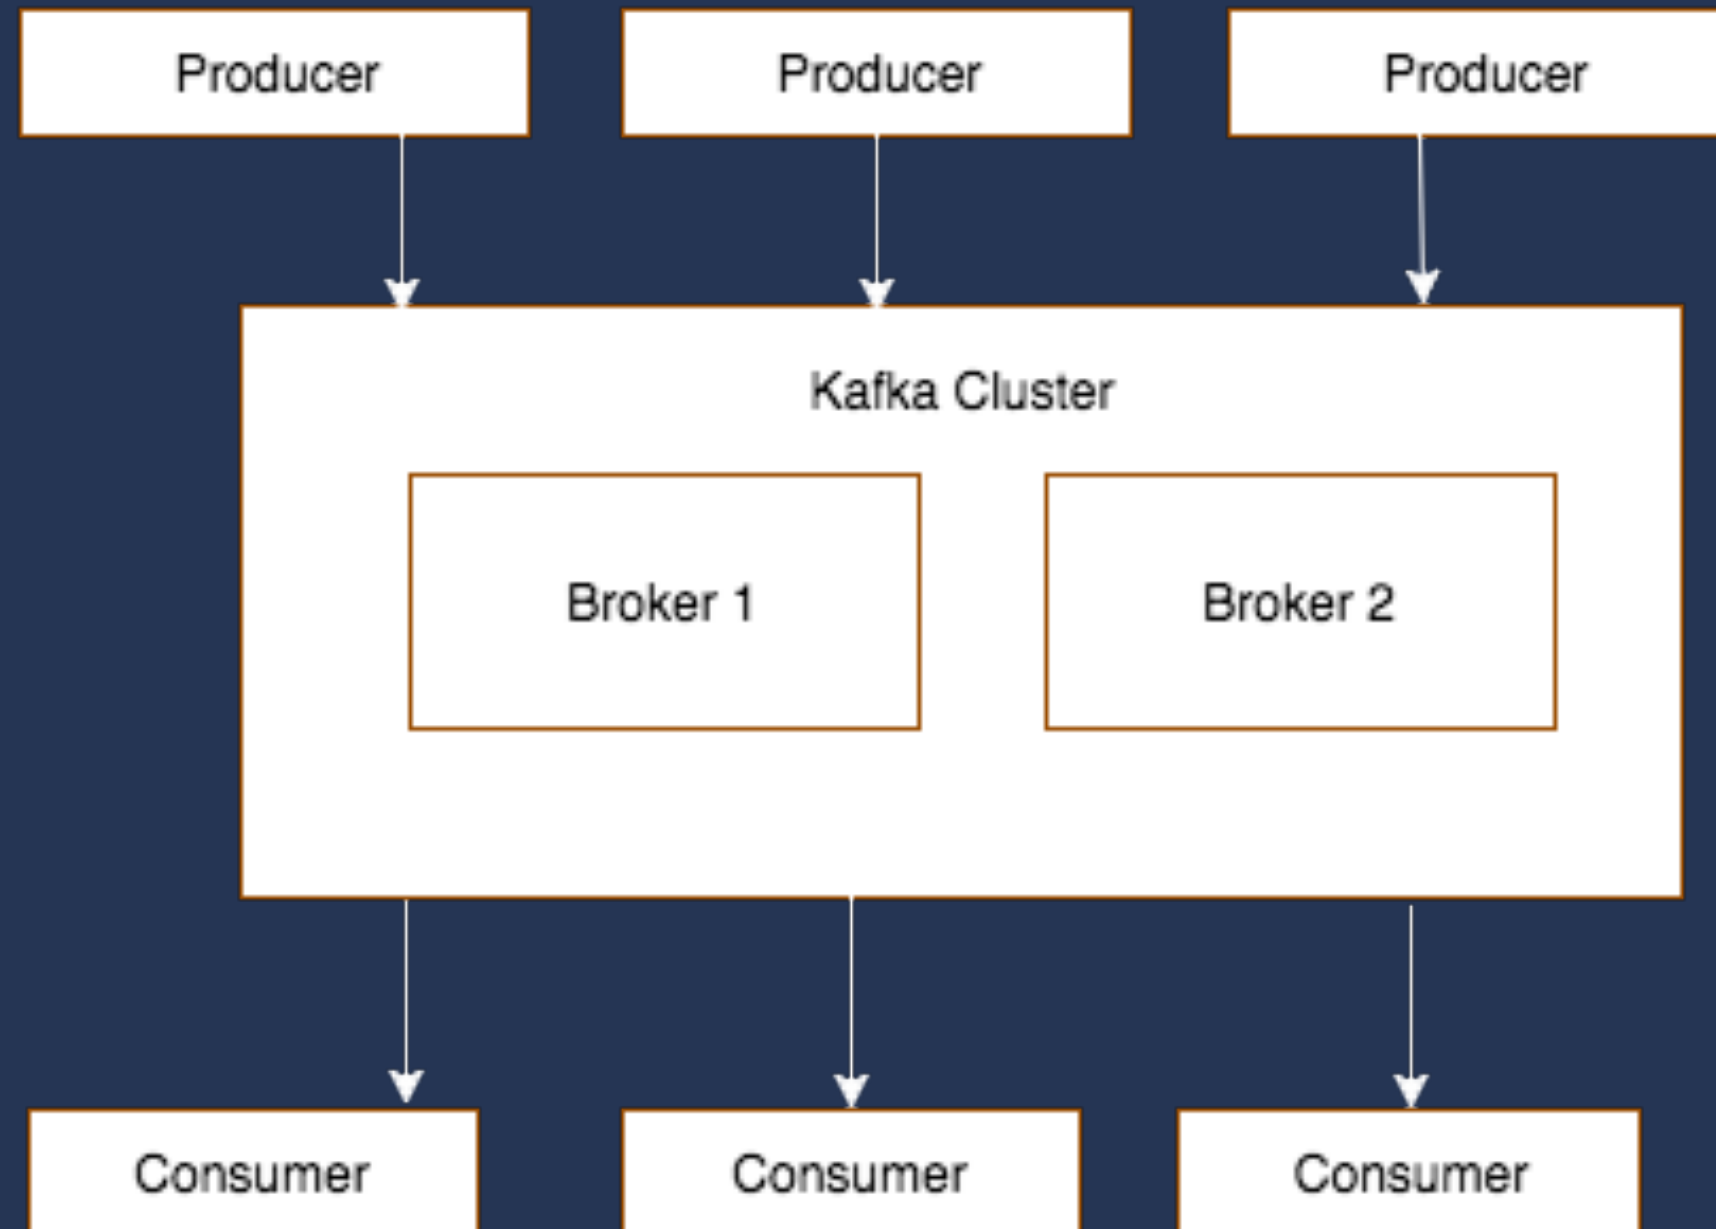


Still, too hard to draw

Sorry, back to kafka

- Pub/Sub messaging rethought as a distributed commit log.
- Messages are divided on different partitioned feeds called *topics*
- Messages are **persistent**, they are not lost when consumed
- Having persistent messages provides a buffer between producers and consumers, if the consumer crashes the data will still be there when it comes back
- It also means you can rewind the log ;)

Anatomy



Topic

Partition



Partition



Consumers & Consumer groups

- Consumers keep track of their offset on the log
- Consumers commit the offset to a special kafka topic to recover in case of a crash
- This gives "at least 1" delivery but you can get "exactly 1"
- Consumer groups automatically distribute partitions among themselves, free scaling (both up and down) up to the number of partitions on the topic

TODO: ADD DEMO

Thank you

albertoavila@gmail.com
@alberteiu